

Assignment2

March 5, 2023

1 Overview

Group 136

Name	BITS ID	Contribution
Vinayak Nayak	2021fc04135	100%
Shreysi Kalra	2021fc04586	100%

Setup Details

- Operating System: Ubuntu 20.04 LTS
- Hadoop: Hadoop 3.2.4
- Python: Python 3.8
- Hive: Hive 3.1.3

Submission Details Summary

- Hadoop is installed and configured in **pseudo-distributed** setup since I have only one node/machine. Hence the default replication factor in the **hdfs-site.xml** is set to 1. In a real-world distributed scenario we would configure it to be 3/5 i.e. odd number > 1 so that we can achieve fault-tolerance and have a majority quorum in case of network-partition or temporary node-failure etc.
- Query 1 is performed using Apache Hive (Which uses Hadoop Map-Reduce as the execution engine beneath it; we haven't configured **tez** or **spark** but use the default **map-reduce** engine for hive)
- Queries 2-7 are performed using Hadoop Map-Reduce with the Hadoop Streaming API for Python.
- All the outputs are stored in the format **/home/query<#>_output/** which are displayed below in chronological order.
- The following videos shows the execution of all our queries with our environment setup and execution engine.
 - [Part 1](#)
 - [Part 2](#)
- The map-reduce code for each of the queries 2-7 is organized in folders with names **query_{#}** respectively.
- The commands for creation of table, loading of data from csv and querying for unique customers from **Germany** as the country of our choice is organized in the folder **query_1**. > Each query is executed using python data analysis libraries like numpy and also subsequently using

Hadoop map-reduce/ Apache hive. The code for both these methods and the corresponding results are also attached in the document below. Both of these results map which ascertain the veracity of the executed queries.

- Configuration files which were modified for utilization of `hadoop` and `hive` are stored in `configuration_files` respectively
- Changes made to `.bashrc` file on my setup are reflected in `bashrc` file attached in `configuration_files`. Please note I had to remove other lines of code from the `.bashrc` file to preserve my privacy.
- There are some considerations which I had to make made in order to account for some anomalous observations in the dataset. I have mentioned them in the `Data Understanding` module below by doing a quick EDA on the same.

```
[ ]: # Have a look at the different output files created as a result of map reduce ↵  
      ↪output  
      !hadoop fs -ls /home/
```

```
Found 7 items  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 19:57  
/home/query2_output  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 19:58  
/home/query3_output  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 19:59  
/home/query4_output  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 20:27  
/home/query5_output  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 20:00  
/home/query6_output  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 20:22  
/home/query7_output  
drwxr-xr-x  - vinayak supergroup          0 2023-03-04 18:33 /home/vinayak
```

```
[ ]: # Have a look at the input csv placed in the hadoop file system  
      !hadoop fs -ls /hdfs_input
```

```
Found 1 items  
-rw-r--r--  1 vinayak supergroup  49722977 2023-03-04 17:46  
/hdfs_input/assignment2_retail_data_utf8.csv
```

```
[ ]: # Have a look at the hive data which has gotten created due to  
      # import of the csv file into hive  
      !hadoop fs -ls /user/hive/warehouse/test.db
```

```
Found 1 items  
drwxr-xr-x  - vinayak supergroup          0 2023-03-05 07:41  
/user/hive/warehouse/test.db/retail
```

2 Data Understanding

```
[ ]: import pandas as pd
df = pd.read_csv(
    "data/Assignment -2 2023 BDS DATA SET online_retail_data.csv",
    encoding="unicode_escape",
)
df.head(2)
```

```
[ ]:      Record No. Invoice StockCode      Description \
0          1  489434      85048  15CM CHRISTMAS GLASS BALL 20 LIGHTS
1          2  489434      79323P      PINK CHERRY LIGHTS

      Quantity      InvoiceDate Price Customer ID      Country
0          12  12-01-2009 07:45   6.95      13085.0  United Kingdom
1          12  12-01-2009 07:45   6.75      13085.0  United Kingdom
```

There are non standard i.e. (not utf-8) characters also in our dataframe. This will cause our map reduce jobs to fail as they wouldn't be rendered in our python map-reduce characters. We will have to convert these into utf-8 values. For eg, consider the following line.

```
[ ]: !sed -ne 31081p data/"Assignment -2 2023 BDS DATA SET online_retail_data.csv"
```

```
31080,491969,gift_0001_80,Dotcomgiftshop Gift Voucher 80.00,1,12/14/2009
17:57:00,69.56,,United Kingdom
```

There is a pound sign in the description which is not getting rendered properly, hence we will have to first convert them into utf-8 characters before passing them to **map-reduce** tasks.

```
[ ]: df.dtypes
```

```
[ ]: Record No.      int64
Invoice            object
StockCode          object
Description         object
Quantity           int64
InvoiceDate        object
Price              float64
Customer ID        float64
Country            object
dtype: object
```

We can see that customer ID has been read as a float. This means there is some issue, let's try to see how many nas are present in our data.

```
[ ]: df.isnull().sum()
```

```
[ ]: Record No.      0
      Invoice         0
      StockCode      0
      Description    2928
      Quantity       0
      InvoiceDate     0
      Price          0
      Customer ID    107927
      Country        0
      dtype: int64
```

Since there are null values in `Customer ID` column, they are all read as nans and the column has been made into a float dtype.

```
[ ]: df[df.Price < 0].head(3)
```

```
[ ]:      Record No.  Invoice StockCode  Description  Quantity \
179403      179404  A506401          B  Adjust bad debt          1
276274      276275  A516228          B  Adjust bad debt          1
403472      403473  A528059          B  Adjust bad debt          1

      InvoiceDate  Price  Customer ID  Country
179403  04/29/2010 13:36:00 -53594.36      NaN  United Kingdom
276274  07/19/2010 11:24:00 -44031.79      NaN  United Kingdom
403472  10/20/2010 12:04:00 -38925.87      NaN  United Kingdom
```

```
[ ]: df[df.Quantity < 0].sample(3)
```

```
[ ]:      Record No.  Invoice StockCode  Description \
17551      17552  C490798    84997D  PINK 3 PIECE MINI DOTS CUTLERY SET
381112      381113  C526110    22470          HEART OF WICKER LARGE
442961      442962  C531557    22271          FELTCRAFT DOLL ROSIE

      Quantity  InvoiceDate  Price  Customer ID  Country
17551         -6  12-08-2009 11:51    3.75    14277.0    France
381112         -1  10-08-2010 13:03    2.55    13113.0  United Kingdom
442961        -144  11-09-2010 10:06    2.55    12454.0    Spain
```

There are negative values in `price` column and `quantity` column which are not actually transactions but some kind of ledger statements from the account or something; we must take this into consideration when using map-reduce on these queries.

Our way of handling

If there is a negative value for quantity or for price, ignore that line item; do not consider it for our analysis.

One more assumption in our analysis is that we're considering the price column as price per unit of the line item mentioned. So our revenue calculations are done after multiplication of quantity

with price and then summing across the line items.

```
[ ]: stock_codes = [
    "ADJUST", "ADJUST2", "AMAZONFEE", "B", "BANK CHARGES", "C2", "C3", "D",
    "DOT", "GIFT", "M", "POST", "S", "TEST001", "TEST002", "m",
]

df[df.StockCode.isin(stock_codes)].drop_duplicates(subset=["StockCode"])[
    ["StockCode", "Description"]].reset_index(drop=True)
```

```
[ ]:
      StockCode      Description
0          POST      POSTAGE
1           D      Discount
2          DOT      DOTCOM POSTAGE
3           M      Manual
4          C2      CARRIAGE
5  BANK CHARGES      Bank Charges
6      TEST001      This is a test product.
7      TEST002      This is a test product.
8      ADJUST  Adjustment by john on 26/01/2010 16
9       GIFT      NaN
10         m      Manual
11         S      SAMPLES
12         B      Adjust bad debt
13      ADJUST2  Adjustment by Peter on Jun 25 2010
14         C3      NaN
15  AMAZONFEE      AMAZON FEE
```

There are some stock codes which are not actually pertinent to any product but it's related to some accounting/finance details. When writing our queries we will need to skip these test products respectively. As seen above these are clearly not items which are sold but services/miscellaneous things which are dumped in the data.

```
[ ]: # Store the dataframe using `utf-8` encoding for our usecase.
df.to_csv("data/assignment2_retail_data_utf8.csv", index=False,
encoding="utf-8")
```

Note

Clarification was sought from Professor Sunil Bhutada regarding the file format conversion and permission was granted to consider `utf-8` conversion as mentioned in the `instructor_clarifications` folder of the conversation with the Professor.

```
[ ]: # For local comparisons, make sure to drop these stock code items
df = df[~df.StockCode.isin(stock_codes)]
```

In the subsequent sections of this document, I have shown the queries and their respective outputs both in Apache ecosystem and in the local ecosystem. The Map-Reduce Job Summary for each of the jobs is attached at the end of the query as requested in the submission requirements respectively.

3 Queries

Query 1 Computation

Total number of unique customers in the “given country”.

We are considering given country as Germany

```
[ ]: df[df.Country == "Germany"]["Customer ID"].nunique()
```

```
[ ]: 68
```

Using Hive

```
# Enter into hive shell
```

```
cd /home/vinayak/apache-hive-3.1.3-bin
```

```
bin/hive
```

```
# Hive commands to create table, load the data and execute the query
```

```
create table test.retail
```

```
(
```

```
record_id int,
```

```
invoice int,
```

```
stockcode string,
```

```
description string,
```

```
quantity int,
```

```
invoicedate string,
```

```
price float,
```

```
customerid int,
```

```
country string
```

```
)
```

```
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde' STORED AS TEXTFILE;
```

```
load data local
```

```
inpath '/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/data/assignment2_retail_data.csv'
```

```
into table test.retail;
```

```
SELECT COUNT(DISTINCT(customerid)) from test.retail WHERE country="Germany";
```

```

hive> use test;
OK
Time taken: 0.016 seconds
hive> SELECT COUNT(DISTINCT(customerid)) from test.retail WHERE country="Germany";
Query ID = vinayak_20230305110502_777b8e7b-e8bb-4eee-9c37-89dfd0d9c4df
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1677932056274_0014, Tracking URL = http://KARZA-HW-0270:8088/proxy/application_1677932056274_0014/
Kill Command = /home/vinayak/hadoop-3.2.4/bin/mapred job -kill job_1677932056274_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-05 11:05:14,868 Stage-1 map = 0%, reduce = 0%
2023-03-05 11:05:21,073 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.62 sec
2023-03-05 11:05:25,164 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.19 sec
MapReduce Total cumulative CPU time: 6 seconds 190 msec
Ended Job = job_1677932056274_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.19 sec HDFS Read: 49733141 HDFS Write: 102 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 190 msec
OK
68
Time taken: 23.429 seconds, Fetched: 1 row(s)
hive> SELECT * from test.retail limit 3;
OK
Record No. Invoice StockCode Description Quantity InvoiceDate Price Customer ID Country
1 489434 85048 15CM CHRISTMAS GLASS BALL 20 LIGHTS 12 12-01-2009 07:45 6.95 13085.0 United Kingdom
2 489434 79323P PINK CHERRY LIGHTS 12 12-01-2009 07:45 6.75 13085.0 United Kingdom
Time taken: 0.131 seconds, Fetched: 3 row(s)
hive>

```

The summary of the corresponding hive query in map-reduce jobs is as follows.

MapReduce Job job_1677932056274_0014

Job Overview

Job Name:	SELECT COUNT(DISTINCT(cu...country="Germany" (Stage-1)
User Name:	vinayak
Queue:	default
State:	SUCCEEDED
Uberized:	false
Submitted:	Sun Mar 05 11:05:05 IST 2023
Started:	Sun Mar 05 11:05:14 IST 2023
Finished:	Sun Mar 05 11:05:24 IST 2023
Elapsed:	10sec
Diagnostics:	
Average Map Time	3sec
Average Shuffle Time	1sec
Average Merge Time	0sec
Average Reduce Time	0sec

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Sun Mar 05 11:05:11 IST 2023	KARZA-HW-0270:8042	logs

Task Summary

Task Type	1	Total	1	Complete
Map	1	1	1	
Reduce	1	1	1	
Attempt Type		Failed	Killed	Successful
Maps	0	0	1	
Reduces	0	0	1	

Query 2 Computation

Country from which the maximum revenue was collected from sales in the month of March 2010.

```
[ ]: df["InvoiceDate"] = pd.to_datetime(df["InvoiceDate"])

[ ]: df["Month_Year"] = (
    df["InvoiceDate"].dt.month.apply(lambda x: str(x))
    + "/"
    + df["InvoiceDate"].dt.year.apply(lambda x: str(x))
)

[ ]: revenue_map = {}
for country, sub_df in df[df["Month_Year"] == "3/2010"].groupby(by="Country"):
    sub_df = sub_df[sub_df.Quantity > 0]
    sub_df = sub_df[sub_df.Price > 0]
    revenue = (sub_df.Quantity * sub_df.Price).sum()
    revenue_map[country] = round(revenue, 2)

[ ]: print(revenue_map)

{'Australia': 429.39, 'Austria': 685.13, 'Bahrain': 548.8, 'Belgium': 629.02,
'Bermuda': 1253.14, 'Channel Islands': 1065.12, 'Cyprus': 2879.19, 'Denmark':
7595.18, 'EIRE': 21778.3, 'France': 8027.59, 'Germany': 15347.08, 'Greece':
522.73, 'Italy': 699.22, 'Japan': 110.4, 'Netherlands': 24241.3, 'Poland':
318.86, 'Portugal': 2399.45, 'Spain': 1456.52, 'Sweden': 1357.58, 'Switzerland':
831.62, 'United Arab Emirates': 1201.52, 'United Kingdom': 670999.26}

[ ]: highest_revenue_country = sorted(
    revenue_map, key=lambda x: revenue_map[x], reverse=True
)[0]
print(highest_revenue_country, revenue_map[highest_revenue_country])
```

United Kingdom 670999.26

Using Hadoop MapReduce

mapper.py

```
#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_2/mapper.py"""
```

```
import sys
from datetime import datetime
```

```
banned_stock_codes = ["ADJUST", "ADJUST2", "AMAZONFEE", "B",
    "BANK CHARGES", "C2", "C3", "D",
    "DOT", "GIFT", "M",
    "POST", "S", "TEST001", "TEST002", "m"]
```



```

banned_stock_codes = set(banned_stock_codes)

# Read the input from standard input
for line in sys.stdin:
    # Remove all the whitespaces
    line = line.strip()

    # Split any given record into corresponding fields
    rec_no, invoice, stock_code, *desc, quantity, invoice_date, price,\
    cust_id, country = line.split(",")

    # We don't want to do any computation for the header of the file,
    # hence skip the first line
    if not country == "Country":

        # Compute the revenue generated
        price = float(price); quantity = int(quantity)
        revenue = price * quantity

        # Parse the date appropriately
        # Some dates are delimited by / and some others are delimited by -
        # Some dates have hour-min-sec in time and some others have hour-min
        if len(invoice_date.split(":")) == 2:
            if "/" in invoice_date:
                dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M")
            else:
                dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M")
        else:
            if "/" in invoice_date:
                dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M:%S")
            else:
                dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M:%S")

        # Only print the dates which belong to the month of March 2010
        if (dt.month == 3) and (dt.year == 2010):
            # Check if the revenue, quantity and prices are positive
            if (price > 0) and (quantity > 0) and (revenue > 0):
                # Check if the given stock code is not in the banned stock codes
                if stock_code not in banned_stock_codes:
                    print(f"{country},{revenue}")

```

reducer.py

```

#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_2/reducer.py"""

```

```

import sys

```

```

from collections import defaultdict
country_revenue_map = defaultdict(lambda: 0)

# Read the input from standard input
for line in sys.stdin:
    # Remove the leading and trailing whitespaces
    line = line.strip()

    # Parse the mapper input
    country, revenue = line.split(',')

    # Sum the revenue into the respective country's account
    country_revenue_map[country] += float(revenue)

# Figure out the country with highest revenue
highest_revenue_country = sorted(country_revenue_map, \
                                  key = lambda x: country_revenue_map[x], \
                                  reverse = True)[0]

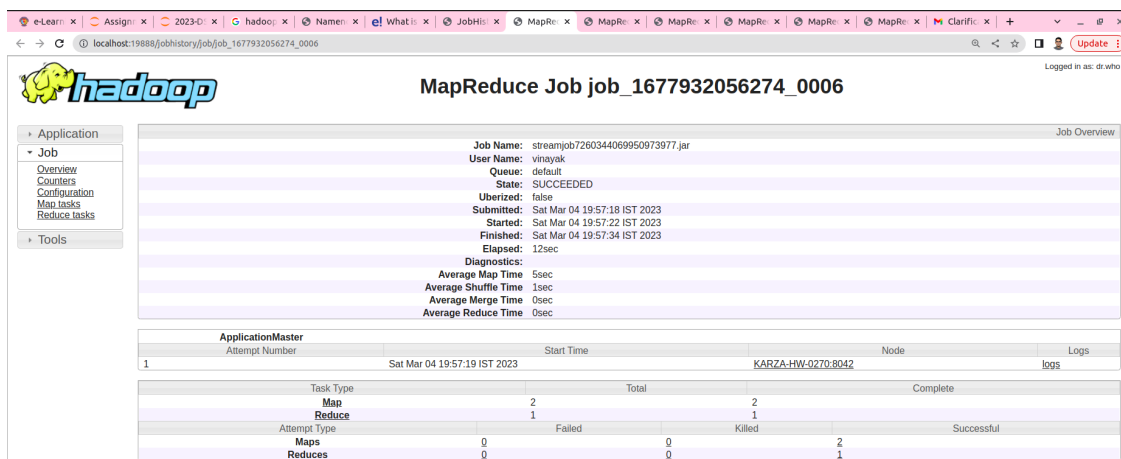
# Print the highest revenue country and the highest revenue to console
print(f"For 03/2010; Highest Revenue Country: {highest_revenue_country}\
Highest Revenue: {country_revenue_map[highest_revenue_country]:.3f}")

hadoop jar /home/vinayak/hadoop-3.2.4/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_2/mapper.py
-mapper mapper.py
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_2/reducer.py
-reducer reducer.py
-input /hdfs_input/assignment2_retail_data_utf8.csv
-output /home/query2_output/

```

```
[ ]: !hdfs dfs -cat /home/query2_output/part-00000
```

For 03/2010; Highest Revenue Country: United Kingdom | Highest Revenue: 670999.261



The screenshot shows the Hadoop MapReduce Job interface for job `job_1677932056274_0006`. The job is in a **SUCCEEDED** state. The user is `vinayak` and the queue is `default`. The job was submitted on `Sat Mar 04 19:57:18 IST 2023` and finished on `Sat Mar 04 19:57:34 IST 2023`, with an elapsed time of `12sec`. The diagnostics show an average map time of `5sec`, average shuffle time of `15sec`, average merge time of `0sec`, and average reduce time of `0sec`.

ApplicationMaster		Start Time	Node	Logs
1	Attempt Number	Sat Mar 04 19:57:19 IST 2023	KARZA-HW-0270:8042	logs

Task Type	2	Total	Complete
Map	2	2	
Reduce	1	1	
Attempt Type	Failed	Killed	Successful
Maps	0	0	2
Reduces	0	0	1

Query 3 Computation

Month of 2010 in which maximum number of items were sold.

```
[ ]: num_items = {}
for my, sub_df in df.groupby(by="Month_Year"):
    sub_df = sub_df[sub_df.Quantity > 0]
    sub_df = sub_df[sub_df.Price > 0]
    sub_df = sub_df[sub_df.InvoiceDate.dt.year == 2010]
    items = sub_df.Quantity.sum()
    num_items[my] = round(items, 2)

[ ]: highest_selling_month = sorted(num_items, key=lambda x: num_items[x],
    ↪reverse=True)[0]
print(
    f"For 2010; Highest Selling Month: {highest_selling_month} | Items Sold: ↪
    ↪{num_items[highest_selling_month]}"
)
```

For 2010; Highest Selling Month: 11/2010 | Items Sold: 727556

Using Hadoop MapReduce

mapper.py

```
#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_3/mapper.py"""

import sys
from datetime import datetime

banned_stock_codes = ["ADJUST", "ADJUST2", "AMAZONFEE", "B",
    "BANK CHARGES", "C2", "C3", "D",
    "DOT", "GIFT", "M",
    "POST", "S", "TEST001", "TEST002", "m"]
banned_stock_codes = set(banned_stock_codes)

# Read the input from standard input
for line in sys.stdin:
    # Remove all the whitespaces
    line = line.strip()

    # Split any given record into corresponding fields
    rec_no, invoice, stock_code, *desc, quantity, invoice_date, \
    price, cust_id, country = line.split(",")

    # We don't want to do any computation for the header of the file,
    # hence skip the first line
    if not country == "Country":
        # Cast the price and quantity appropriately
```

```

price = float(price); quantity = int(quantity)

# Parse the date appropriately
# Some dates are delimited by / and some others are delimited by -
# Some dates have hour-min-sec in time and some others have hour-min
if len(invoice_date.split(":")) == 2:
    if "/" in invoice_date:
        dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M")
    else:
        dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M")
else:
    if "/" in invoice_date:
        dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M:%S")
    else:
        dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M:%S")

# Only print the those transactions which happened in the year 2010
if dt.year == 2010:
    # Check if the quantity and prices are positive
    if (price > 0) and (quantity > 0):
        # Check if given stock code is not in banned stock codes
        if not stock_code in banned_stock_codes:
            print(f"{dt.month},{quantity}")

```

reducer.py

```

#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_3/reducer.py"""

import sys
from collections import defaultdict
month_quantity_map = defaultdict(lambda: 0)

# Read the input from standard input
for line in sys.stdin:
    # Remove the leading and trailing whitespaces
    line = line.strip()

    # Parse the mapper input
    month, quantity = line.split(',')

    # Sum the quantity into the respective month's account
    month_quantity_map[month] += int(quantity)

# Figure out the month with highest items sold
highest_selling_month = sorted(month_quantity_map, key = lambda x: month_quantity_map[x], \
                                reverse = True)[0]

```

```
# Print the highest revenue country and the highest revenue to console
print(f"For 2010; Highest Selling Month: {highest_selling_month:0<2} | \
Items Sold: {month_quantity_map[highest_selling_month]}")
```

```
hadoop jar /home/vinayak/hadoop-3.2.4/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_3/mapper.py
-mapper mapper.py
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_3/reducer.py
-reducer reducer.py
-input /hdfs_input/assignment2_retail_data_utf8.csv
-output /home/query3_output/
```

```
[ ]: !hdfs dfs -cat /home/query3_output/part-00000
```

For 2010; Highest Selling Month: 11 | Items Sold: 727556

The screenshot shows the Hadoop MapReduce Job interface for job_1677932056274_0007. The job is in a 'SUCCEEDED' state. The interface includes a sidebar with navigation options like 'Application', 'Job', 'Overview', 'Counters', 'Configuration', 'Map tasks', and 'Reduce tasks'. The main content area displays job details such as Job Name, User Name, Queue, State, Submitted, Started, Finished, Elapsed, and Diagnostics. Below this, there is a table for 'ApplicationMaster' showing attempt numbers, start times, and node information. At the bottom, a table shows the progress of 'Map' and 'Reduce' tasks, including counts for 'Failed', 'Killed', and 'Successful' attempts.

Task Type	2	Total	2	Complete
Map	2		2	
Reduce	1		1	
Attempt Type		Failed	Killed	Successful
Maps	0	0	2	
Reduces	0	0	1	

Query 4 Computation

In the month of January 2010, find the country in which maximum number of items were sold

```
[ ]: country_items = {}
for country, sub_df in df[df["Month_Year"] == "1/2010"].groupby(by="Country"):
    sub_df = sub_df[sub_df.Quantity > 0]
    sub_df = sub_df[sub_df.Price > 0]
    items = sub_df.Quantity.sum()
    country_items[country] = items
```

```
[ ]: highest_selling_country = sorted(
    country_items, key=lambda x: country_items[x], reverse=True
)[0]
print(
    f"For 01/2010; Country selling max items: {highest_selling_country} | ␣
↪Volume Sold: {country_items[highest_selling_country]}"
)
```

For 01/2010; Country selling max items: United Kingdom | Volume Sold: 257473

Using Hadoop MapReduce

mapper.py

```
#!/home/vinayak/anaconda3/bin/python
```

```
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_4/mapper.py"""
```

```
import sys
```

```
from datetime import datetime
```

```
banned_stock_codes = ["ADJUST", "ADJUST2", "AMAZONFEE", "B",
                      "BANK CHARGES", "C2", "C3", "D",
                      "DOT", "GIFT", "M",
                      "POST", "S", "TEST001", "TEST002", "m"]
```

```
banned_stock_codes = set(banned_stock_codes)
```

```
# Read the input from standard input
```

```
for line in sys.stdin:
```

```
# Remove all the whitespaces
```

```
line = line.strip()
```

```
# Split any given record into corresponding fields
```

```
rec_no, invoice, stock_code, *desc, quantity, invoice_date, \
price, cust_id, country = line.split(",")
```

```
# We don't want to do any computation for the header of the file,
```

```
# hence skip the first line
```

```
if not country == "Country":
```

```

price = float(price); quantity = int(quantity)

# Parse the date appropriately
# Some dates are delimited by / and some others are delimited by -
# Some dates have hour-min-sec in time and some others have hour-min
if len(invoice_date.split(":")) == 2:
    if "/" in invoice_date:
        dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M")
    else:
        dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M")
else:
    if "/" in invoice_date:
        dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M:%S")
    else:
        dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M:%S")

# Only select the dates which belong to the month of January 2010
if (dt.month == 1) and (dt.year == 2010):
    # Check if the quantity and prices are positive
    if (price > 0) and (quantity > 0):
        # Check if the stock code is not in the banned stock codes
        if stock_code not in banned_stock_codes:
            print(f"{country},{quantity}")

```

reducer.py

```

#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_4/reducer.py"""

import sys
from collections import defaultdict
country_quantity_map = defaultdict(lambda: 0)

# Read the input from standard input
for line in sys.stdin:
    # Remove the leading and trailing whitespaces
    line = line.strip()

    # Parse the mapper input
    country, quantity = line.split(',')

    # Sum the revenue into the respective country's account
    country_quantity_map[country] += int(quantity)

# Figure out the country with highest quantity of items sold
highest_selling_country = sorted(country_quantity_map, key = lambda x: country_quantity_map[x])

```

```
reverse = True)[0]
```

```
print(f"For 01/2010; Country selling max items: {highest_selling_country}\  
| Volume Sold: {country_quantity_map[highest_selling_country]}")
```

```
hadoop jar /home/vinayak/hadoop-3.2.4/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar  
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_4/mapper.py  
-mapper mapper.py  
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_4/reducer.py  
-reducer reducer.py  
-input /hdfs_input/assignment2_retail_data_utf8.csv  
-output /home/query4_output/
```

```
[ ]: !hdfs dfs -cat /home/query4_output/part-00000
```

For 01/2010; Country selling max items: United Kingdom | Volume Sold: 257473

MapReduce Job job_1677932056274_0008

Job Overview

Job Name: streamjob5225949446050446358.jar
User Name: vinayak
Queue: default
State: SUCCEEDED
Uberized: false
Submitted: Sat Mar 04 19:59:00 IST 2023
Started: Sat Mar 04 19:59:04 IST 2023
Finished: Sat Mar 04 19:59:16 IST 2023
Elapsed: 12sec
Diagnostics:
Average Map Time: 5sec
Average Shuffle Time: 1sec
Average Merge Time: 0sec
Average Reduce Time: 0sec

ApplicationMaster

Attempt Number	Start Time	Node	Logs
1	Sat Mar 04 19:59:01 IST 2023	KARZA-HW-0270:8042	logs

Task Summary

Task Type	2	Total	2
Map	2	2	Complete
Reduce	1	1	1
Attempt Type	Failed	Killed	Successful
Maps	0	0	2
Reduces	0	0	1

Query 5 Computation

The StockCode of the item with the highest number of sales in the “given country” in the year 2010

We are considering given country as Germany

```
[ ]: stock_code_sales = {}
for stock_code, sub_df in df[
    (df.InvoiceDate.dt.year == 2010) & (df.Country == "Germany")
].groupby(by="StockCode"):
    sub_df = sub_df[sub_df.Quantity > 0]
    sub_df = sub_df[sub_df.Price > 0]
    sales = sub_df.Quantity.sum()
    stock_code_sales[stock_code] = sales

[ ]: highest_sale_item = sorted(
    stock_code_sales, key=lambda x: stock_code_sales[x], reverse=True
)[0]
print(
    f"For Germany; Stock code with highest sale: {highest_sale_item} | Volume_
    ↪Sold: {stock_code_sales[highest_sale_item]}"
)
```

For Germany; Stock code with highest sale: 22326 | Volume Sold: 1543

Using Hadoop MapReduce

mapper.py

```
#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_5/mapper.py"""

import sys
from datetime import datetime

banned_stock_codes = ["ADJUST", "ADJUST2", "AMAZONFEE", "B",
    "BANK CHARGES", "C2", "C3", "D",
    "DOT", "GIFT", "M",
    "POST", "S", "TEST001", "TEST002", "m"]
banned_stock_codes = set(banned_stock_codes)

# Read the input from standard input
for line in sys.stdin:
    # Remove all the whitespaces
    line = line.strip()

    # Split any given record into corresponding fields
    rec_no, invoice, stock_code, *desc, quantity, invoice_date, price, \
    cust_id, country = line.split(",")
```

```

# We only want to consider Germany's sales
if country == "Germany":

    price = float(price); quantity = int(quantity)

    # Parse the date appropriately
    # Some dates are delimited by / and some others are delimited by -
    # Some dates have hour-min-sec in time and some others have hour-min
    if len(invoice_date.split(":")) == 2:
        if "/" in invoice_date:
            dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M")
        else:
            dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M")
    else:
        if "/" in invoice_date:
            dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M:%S")
        else:
            dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M:%S")

    # Only select the dates which belong to the month of December 2010
    if (dt.year == 2010):
        # Check if the quantity and prices are positive
        if (price > 0) and (quantity > 0):
            # Check if the stock code is not in the banned stock codes
            if stock_code not in banned_stock_codes:
                print(f"{stock_code},{quantity}")

```

reducer.py

```

#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_5/reducer.py"""

import sys
from collections import defaultdict
stock_code_sales = defaultdict(lambda: 0)

# Read the input from standard input
for line in sys.stdin:
    # Remove the leading and trailing whitespaces
    line = line.strip()

    # Parse the mapper input
    stock_code, quantity = line.split(',')

    # Sum the revenue into the respective country's account
    stock_code_sales[stock_code] += int(quantity)

```

```
# Figure out the country with highest quantity of items sold
```

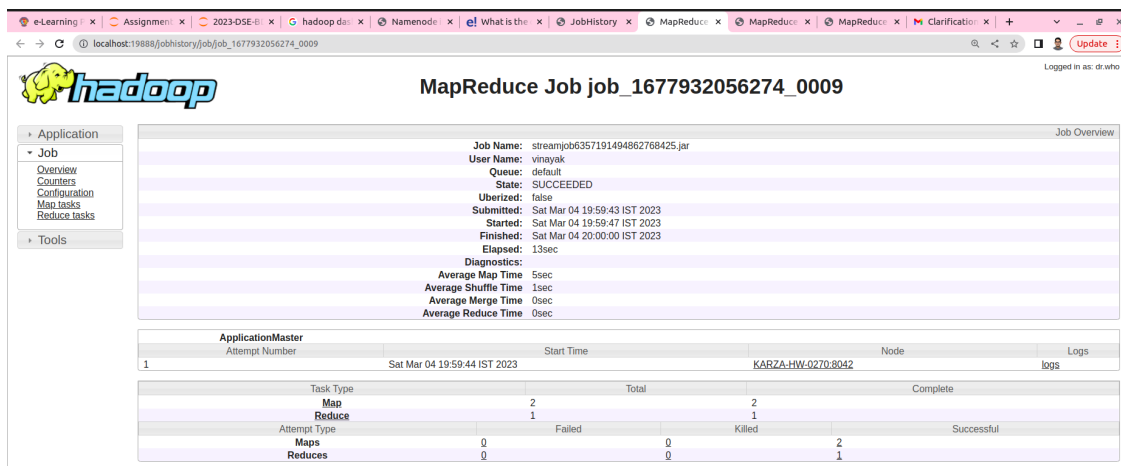
```
highest_sale_item = sorted(stock_code_sales, key = lambda x: stock_code_sales[x], \
                           reverse = True)[0]
```

```
print(f"For Germany; Stock Code giving highest sale: {highest_sale_item} \
| Volume Sold: {stock_code_sales[highest_sale_item]:.2f}")
```

```
hadoop jar /home/vinayak/hadoop-3.2.4/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_5/mapper.py
-mapper mapper.py
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_5/reducer.py
-reducer reducer.py
-input /hdfs_input/assignment2_retail_data_utf8.csv
-output /home/query5_output/
```

```
[ ]: !hdfs dfs -cat /home/query5_output/part-00000
```

For Germany; Stock Code giving highest sale: 22326 | Volume Sold: 1543.00



The screenshot displays the Hadoop MapReduce Job interface for job_id 1677932056274_0009. The job is in a 'SUCCEEDED' state. The interface includes a sidebar with navigation options like 'Application', 'Job', 'Overview', 'Counters', 'Configuration', 'Map tasks', and 'Reduce tasks'. The main content area shows job details such as Job Name, User Name, Queue, State, Uberized, Submitted, Started, Finished, Elapsed, and Diagnostics. Below this, there is a table for 'ApplicationMaster' showing attempt numbers, start times, and nodes. At the bottom, a table provides a summary of task types (Map, Reduce) and their progress (Total, Complete, Failed, Killed, Successful).

Task Type	Total	Complete
Map	2	2
Reduce	1	1

Attempt Type	Failed	Killed	Successful
Maps	0	0	2
Reduces	0	0	1

Query 6 Computation

StockCode of the item for which the maximum revenue was received by sales in the month of December 2010.

```
[ ]: stock_revenue_map = {}
for stock_code, sub_df in df[df["Month_Year"] == "12/2010"].
    ↪groupby(by="StockCode"):
    sub_df = sub_df[sub_df.Quantity > 0]
    sub_df = sub_df[sub_df.Price > 0]
    revenue = (sub_df.Quantity * sub_df.Price).sum()
    stock_revenue_map[stock_code] = round(revenue, 2)

[ ]: highest_revenue_item = sorted(
    stock_revenue_map, key=lambda x: stock_revenue_map[x], reverse=True
)[0]
print(
    f"For 12/2010; Stock Code giving highest revenue: {highest_revenue_item} |
    ↪Revenue Accrued: {stock_revenue_map[highest_revenue_item]}"
)
```

For 12/2010; Stock Code giving highest revenue: 22423 | Revenue Accrued: 13478.98

Using Hadoop MapReduce

mapper.py

```
#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_6/mapper.py"""
```

```
import sys
from datetime import datetime

banned_stock_codes = ["ADJUST", "ADJUST2", "AMAZONFEE", "B",
                      "BANK CHARGES", "C2", "C3", "D",
                      "DOT", "GIFT", "M",
                      "POST", "S", "TEST001", "TEST002", "m"]
banned_stock_codes = set(banned_stock_codes)

# Read the input from standard input
for line in sys.stdin:
    # Remove all the whitespaces
    line = line.strip()

    # Split any given record into corresponding fields
    rec_no, invoice, stock_code, *desc, quantity, invoice_date, price, \
    cust_id, country = line.split(",")

    # We don't want to do any computation for the header of the file,
```

```

# hence skip the first line
if not country == "Country":

    price = float(price); quantity = int(quantity)
    revenue = price * quantity

    # Parse the date appropriately
    # Some dates are delimited by / and some others are delimited by -
    # Some dates have hour-min-sec in time and some others have hour-min
    if len(invoice_date.split(":")) == 2:
        if "/" in invoice_date:
            dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M")
        else:
            dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M")
    else:
        if "/" in invoice_date:
            dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M:%S")
        else:
            dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M:%S")

    # Only select the dates which belong to the month of December 2010
    if (dt.month == 12) and (dt.year == 2010):
        # Check if the quantity and prices are positive
        if (price > 0) and (quantity > 0) and (revenue > 0):
            # Check if the stock code is not in the banned stock codes
            if stock_code not in banned_stock_codes:
                print(f"{stock_code},{revenue}")

```

reducer.py

```

#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_6/reducer.py"""

import sys
from collections import defaultdict
stock_revenue_map = defaultdict(lambda: 0)

# Read the input from standard input
for line in sys.stdin:
    # Remove the leading and trailing whitespaces
    line = line.strip()

    # Parse the mapper input
    stock_code, revenue = line.split(',')

    # Sum the revenue into the respective country's account
    stock_revenue_map[stock_code] += float(revenue)

```

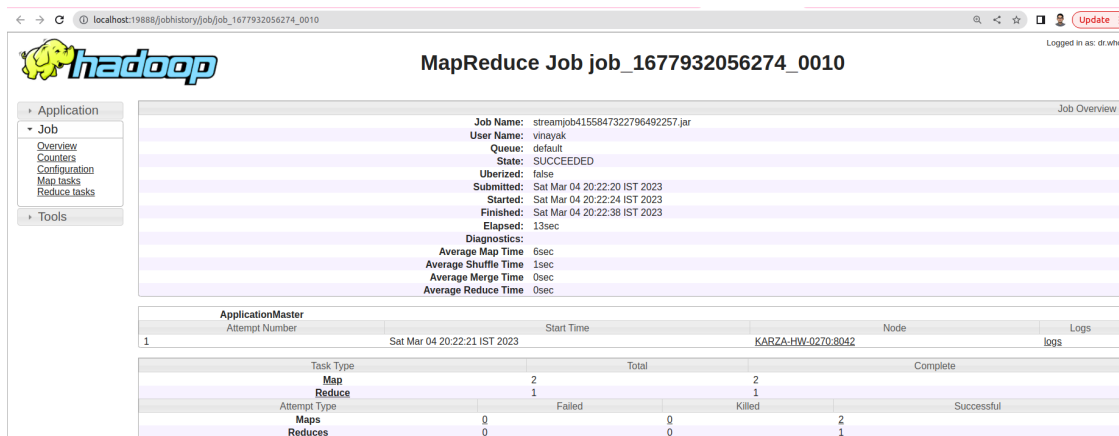
```
# Figure out the country with highest quantity of items sold
highest_revenue_item = sorted(stock_revenue_map, key = lambda x: stock_revenue_map[x], \
                               reverse = True)[0]
```

```
print(f"For 12/2010; Stock Code giving highest revenue: {highest_revenue_item} \
| Revenue Accrued: {stock_revenue_map[highest_revenue_item]:.2f}")
```

```
hadoop jar /home/vinayak/hadoop-3.2.4/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_6/mapper.py
-mapper mapper.py
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_6/reducer.py
-reducer reducer.py
-input /hdfs_input/assignment2_retail_data_utf8.csv
-output /home/query6_output/
```

```
[ ]: !hdfs dfs -cat /home/query6_output/part-00000
```

For 12/2010; Stock Code giving highest revenue: 22423 | Revenue Accrued: 13478.98



The screenshot shows the Hadoop MapReduce Job interface for job_id_1677932056274_0010. The job is in a 'SUCCEEDED' state. The interface includes a sidebar with navigation options like 'Application', 'Job', 'Overview', 'Counters', 'Configuration', 'Main tasks', and 'Reduce tasks'. The main content area displays job details such as Job Name, User Name, Queue, State, and various timing metrics. Below this, there is a table showing the progress of the job, including the number of tasks completed, failed, and killed.

Task Type	2	Total	2	Complete
Map	1	1	1	1
Reduce	1	1	1	1
Attempt Type	Failed	Killed	Successful	
Maps	0	0	2	
Reduces	0	0	1	

Query 7 Computation

The country in which minimum number of sales happened in 2010.

```
[ ]: country_sales = {}
for country, sub_df in df[df.InvoiceDate.dt.year == 2010].groupby(by="Country"):
    sub_df = sub_df[sub_df.Quantity > 0]
    sub_df = sub_df[sub_df.Price > 0]
    sales = sub_df.Quantity.sum()
    country_sales[country] = round(sales, 2)

[ ]: lowest_2010_sales_country = sorted(
    country_sales, key=lambda x: country_sales[x], reverse=False
)[0]
print(
    f"For 2010, Country with lowest quantity sales: {lowest_2010_sales_country} |
    ↪ Total items sold: {country_sales[lowest_2010_sales_country]}"
)
```

For 2010, Country with lowest quantity sales: Lebanon | Total items sold: 72

Using Hadoop MapReduce

mapper.py

```
#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_7/mapper.py"""

import sys
from datetime import datetime

banned_stock_codes = ["ADJUST", "ADJUST2", "AMAZONFEE", "B",
                      "BANK CHARGES", "C2", "C3", "D",
                      "DOT", "GIFT", "M",
                      "POST", "S", "TEST001", "TEST002", "m"]
banned_stock_codes = set(banned_stock_codes)

# Read the input from standard input
for line in sys.stdin:
    # Remove all the whitespaces
    line = line.strip()

    # Split any given record into corresponding fields
    rec_no, invoice, stock_code, *desc, quantity, invoice_date, price, \
    cust_id, country = line.split(",")

    # We don't want to do any computation for the header of the file,
    # hence skip the first line
    if not country == "Country":
```

```

price = float(price); quantity = int(quantity)
revenue = price * quantity

# Parse the date appropriately
# Some dates are delimited by / and some others are delimited by -
# Some dates have hour-min-sec in time and some others have hour-min
if len(invoice_date.split(":")) == 2:
    if "/" in invoice_date:
        dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M")
    else:
        dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M")
else:
    if "/" in invoice_date:
        dt = datetime.strptime(invoice_date, "%m/%d/%Y %H:%M:%S")
    else:
        dt = datetime.strptime(invoice_date, "%m-%d-%Y %H:%M:%S")

# Only select the dates which belong to the year 2010
if dt.year == 2010:
    # Check if the quantity and prices are positive
    if (price > 0) and (quantity > 0):
        # Check if the stock code is not in the banned stock codes
        if stock_code not in banned_stock_codes:
            print(f"{country},{quantity}")

```

reducer.py

```

#!/home/vinayak/anaconda3/bin/python
"""/home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_7/reducer.py"""

import sys
from collections import defaultdict
country_sales_map = defaultdict(lambda: 0)

# Read the input from standard input
for line in sys.stdin:
    # Remove the leading and trailing whitespaces
    line = line.strip()

    # Parse the mapper input
    country, quantity = line.split(',')

    # Sum the revenue into the respective country's account
    country_sales_map[country] += int(quantity)

# Figure out the country with highest quantity of items sold
lowest_sales_item = sorted(country_sales_map, key = lambda x: country_sales_map[x],\
                             reverse = False)[0]

```



```
print(f"For the year 2010; Country with lowest quantity of sales: {lowest_sales_item} \
| Total items sold: {country_sales_map[lowest_sales_item]:.2f}")
```

Execute the hadoop map-reduce job

```
hadoop jar /home/vinayak/hadoop-3.2.4/share/hadoop/tools/lib/hadoop-streaming-3.2.4.jar
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_7/mapper.py
-mapper mapper.py
-file /home/vinayak/Desktop/PGM_BDS_Assignment/BDS_Assignment_2/query_7/reducer.py
-reducer reducer.py
-input /hdfs_input/assignment2_retail_data_utf8.csv
-output /home/query7_output/
```

```
[ ]: !hdfs dfs -cat /home/query7_output/part-00000
```

For the year 2010; Country with lowest quantity of sales: Lebanon | Total items sold: 72.00

The screenshot displays the Hadoop MapReduce Job interface for job `job_1677932056274_0011`. The job is in a **SUCCEEDED** state. The interface includes a sidebar with navigation options like Application, Job, Overview, Counters, Configuration, Map tasks, and Reduce tasks. The main content area shows job details such as Job Name, User Name, Queue, State, Uberized, Submitted, Started, Finished, Elapsed, and Diagnostics. Below this, there is a table for ApplicationMaster attempts and a detailed task progress table.

Task Type	Map	Reduce	Total	Complete
Attempt Type	0	1	1	1
Maps	0	0	0	2
Reduces	0	0	0	1